

Listing Maximal k -Plexes in Large Real-World Graphs

Zhengren Wang, Yi Zhou, Mingyu Xiao, Bakhadyr Khoussainov

Zhengren Wang

Algorithms and Logic Group in UESTC

April 11, 2022



- 1 Background
- 2 Algorithms
- 3 Implementation Techniques
- 4 Experiments

1 Background

2 Algorithms

3 Implementation Techniques

4 Experiments

Finding Cohesive Groups

Finding ***cohesive groups*** (or ***communities***) has received attention from various areas.

- In the WWW, identify clients sharing similar interests and serve them with a common proxy to reduce network traffic.
- In social networks, discover closely related individuals.
- In biological networks, predict the structure and function of protein.



Clique Model

Naturally, cohesive groups can be modeled with **Cliques**.
A clique is a subgraph where vertices are pairwise connected, i.e., a complete graph.




Figure 1: K_7




Figure 2: K_{12}

k -Plex Model

Due to various reasons like *data noise*, communities rarely appear as cliques.
 k -Plexes allow every vertex missing at most $k - 1$ links to other vertices.




Figure 3: 1, 2, 3, 4-plex

Properties

Lemma 1 (Hereditary Property)

- Any induced subgraph of a k -plex is still a k -plex.
- A k -plex is maximal if it is not a subgraph of any larger k -plex.

Lemma 2 (Distance Property)

- Any k -plex with at least $2k - 1$ vertices has its diameter at most 2.
- A k -plex with at most $2k - 2$ vertices may be unconnected.




Figure 4: unconnected 2, 3, 4, 5-plex

Translated Problems

We model cohesive groups with **maximal k -plexes**.

Problem 1 (Listing maximal k -plexes)

Given a graph G , a positive integer k , list all maximal k -plexes from G .

Problem 2 (Listing large maximal k -plexes)

Given a graph G , two positive integers k and l where $l \geq 2k - 1$, list all maximal k -plexes with at least l vertices.

1 Background

2 Algorithms

3 Implementation Techniques

4 Experiments

The Bron-Kerbosch Algorithm

Our algorithm stems from the Bron-Kerbosch algorithm, say *BKPLEX*.

BKPLEX accepts three sets P , C and X as parameters,

- P : **plex vertices** of the growing k -plex,
- C : **candidate vertices** for further branching,
- X : **excluded vertices** to avoid non-maximal solutions.

then lists maximal k -plexes $G[P']$ with three properties:

- $P \subseteq P'$
- $P' \subseteq P \cup C$
- $\forall v \in X$, the subgraph $G[\{v\} \cup P']$ is not a k -plex.

Hereditary Prop. Any induced subgraph of a k -plex is still a k -plex.

The Bron-Kerbosch Algorithm

BKPLEX branches by doing bipartition recursively.¹




Figure 5: An example of BKPLEX.

¹Simplified for clarity. In fact, a variant of bipartition.

Degeneracy Ordering

- $\eta = v_1, \dots, v_n$ is called *degeneracy ordering* (*core ordering*) if each vertex v_i has the minimum degree in the induced subgraph $G[\{v_i, \dots, v_n\}]$.
- Given a degeneracy ordering $\eta = v_1, \dots, v_n$, the degree of v_i in $G[\{v_i, \dots, v_n\}]$ is called the *core number* of v_i .
- For any degeneracy ordering of the same graph, the largest core number among all vertices is a constant D called *degeneracy*.
- Due to the sparsity of many real-world graphs, $D \ll \Delta \ll n$ where Δ is the maximum degree.




Figure 6: An example of degeneracy ordering.

Degeneracy Ordering

- $\eta = v_1, \dots, v_n$ is called *degeneracy ordering* (*core ordering*) if each vertex v_i has the minimum degree in the induced subgraph $G[\{v_i, \dots, v_n\}]$.
- Given a degeneracy ordering $\eta = v_1, \dots, v_n$, the degree of v_i in $G[\{v_i, \dots, v_n\}]$ is called the *core number* of v_i .
- For any degeneracy ordering of the same graph, the largest core number among all vertices is a constant D called *degeneracy*.
- Due to the sparsity of many real-world graphs, $D \ll \Delta \ll n$ where Δ is the maximum degree.




Figure 6: An example of degeneracy ordering.

-

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡

Degeneracy Ordering

- $\eta = v_1, \dots, v_n$ is called *degeneracy ordering* (*core ordering*) if each vertex v_i has the minimum degree in the induced subgraph $G[\{v_i, \dots, v_n\}]$.
- Given a degeneracy ordering $\eta = v_1, \dots, v_n$, the degree of v_i in $G[\{v_i, \dots, v_n\}]$ is called the *core number* of v_i .
- For any degeneracy ordering of the same graph, the largest core number among all vertices is a constant D called *degeneracy*.
- Due to the sparsity of many real-world graphs, $D \ll \Delta \ll n$ where Δ is the maximum degree.




Figure 6: An example of degeneracy ordering.

Degeneracy Ordering

- $\eta = v_1, \dots, v_n$ is called *degeneracy ordering* (*core ordering*) if each vertex v_i has the minimum degree in the induced subgraph $G[\{v_i, \dots, v_n\}]$.
- Given a degeneracy ordering $\eta = v_1, \dots, v_n$, the degree of v_i in $G[\{v_i, \dots, v_n\}]$ is called the *core number* of v_i .
- For any degeneracy ordering of the same graph, the largest core number among all vertices is a constant D called *degeneracy*.
- Due to the sparsity of many real-world graphs, $D \ll \Delta \ll n$ where Δ is the maximum degree.




Figure 6: An example of degeneracy ordering.

Degeneracy Ordering

- $\eta = v_1, \dots, v_n$ is called *degeneracy ordering* (*core ordering*) if each vertex v_i has the minimum degree in the induced subgraph $G[\{v_i, \dots, v_n\}]$.
- Given a degeneracy ordering $\eta = v_1, \dots, v_n$, the degree of v_i in $G[\{v_i, \dots, v_n\}]$ is called the *core number* of v_i .
- For any degeneracy ordering of the same graph, the largest core number among all vertices is a constant D called *degeneracy*.
- Due to the sparsity of many real-world graphs, $D \ll \Delta \ll n$ where Δ is the maximum degree.




Figure 6: An example of degeneracy ordering.

Degeneracy Ordering

- $\eta = v_1, \dots, v_n$ is called *degeneracy ordering* (*core ordering*) if each vertex v_i has the minimum degree in the induced subgraph $G[\{v_i, \dots, v_n\}]$.
- Given a degeneracy ordering $\eta = v_1, \dots, v_n$, the degree of v_i in $G[\{v_i, \dots, v_n\}]$ is called the *core number* of v_i .
- For any degeneracy ordering of the same graph, the largest core number among all vertices is a constant D called *degeneracy*.
- Due to the sparsity of many real-world graphs, $D \ll \Delta \ll n$ where Δ is the maximum degree.




Figure 6: An example of degeneracy ordering.

Degeneracy Ordering

- $\eta = v_1, \dots, v_n$ is called *degeneracy ordering* (*core ordering*) if each vertex v_i has the minimum degree in the induced subgraph $G[\{v_i, \dots, v_n\}]$.
- Given a degeneracy ordering $\eta = v_1, \dots, v_n$, the degree of v_i in $G[\{v_i, \dots, v_n\}]$ is called the *core number* of v_i .
- For any degeneracy ordering of the same graph, the largest core number among all vertices is a constant D called *degeneracy*.
- Due to the sparsity of many real-world graphs, $D \ll \Delta \ll n$ where Δ is the maximum degree.




Figure 6: An example of degeneracy ordering.

Degeneracy Ordering

- $\eta = v_1, \dots, v_n$ is called *degeneracy ordering* (*core ordering*) if each vertex v_i has the minimum degree in the induced subgraph $G[\{v_i, \dots, v_n\}]$.
- Given a degeneracy ordering $\eta = v_1, \dots, v_n$, the degree of v_i in $G[\{v_i, \dots, v_n\}]$ is called the *core number* of v_i .
- For any degeneracy ordering of the same graph, the largest core number among all vertices is a constant D called *degeneracy*.
- Due to the sparsity of many real-world graphs, $D \ll \Delta \ll n$ where Δ is the maximum degree.




Figure 6: An example of degeneracy ordering.

Recent Algorithms³

Pivot Heuristic Zhou et al. (2020) proposed *BKPivot* with a *pivot* heuristic in the branch scheme of BKPLEX, and improved its running time from $O^*(2^n)$ to $O^*(\gamma_k^n)$, where $\gamma_k < 2$.²

Graph Decomposition Conte et al. (2018) proposed *D2K*, a decomposition-based algorithm for listing k -plexes with the diameter at most 2. For each vertex, *D2K* builds a local subgraph and then adopts BKPLEX to list maximal k -plexes locally which runs in $O^*(2^{D\Delta})$.

Distance Prop. Any k -plex with at least $2k - 1$ vertices has the diameter at most 2.

²The notation O^* omits the polynomial factors.

³Simplified for clarity.

Recent Algorithms³

Pivot Heuristic Zhou et al. (2020) proposed *BKPivot* with a *pivot* heuristic in the branch scheme of BKPLEX, and improved its running time from $O^*(2^n)$ to $O^*(\gamma_k^n)$, where $\gamma_k < 2$.²

Graph Decomposition Conte et al. (2018) proposed *D2K*, a decomposition-based algorithm for listing k -plexes with the diameter at most 2. For each vertex, D2K builds a local subgraph and then adopts BKPLEX to list maximal k -plexes locally which runs in $O^*(2^{D\Delta})$.

Distance Prop. Any k -plex with at least $2k - 1$ vertices has the diameter at most 2.

We combine them and propose *ListPlex* running in $O^*(\gamma_k^D)$.

²The notation O^* omits the polynomial factors.

³Simplified for clarity.

ListPlex

Listing All Maximal k -Plexes

Based on Distance Property, ListPlex divides its task into two parts.

- *Part I*: k -plexes with size at most $2k - 2$. (Solved directly by BKPLex)
- *Part II*: k -plexes with size at least $2k - 1$.

Distance Prop. Any k -plex with at least $2k - 1$ vertices has the diameter at most 2.

ListPlex

Part II

Procedure

- Sort V by a degeneracy ordering $\eta = v_1 \dots v_n$.
 - $N^k(v)$ denotes k -hop neighbors of v .
 - $N_{\succ}^k(v_i)$ denotes $N^k(v_i) \cap \{v_{i+1}, \dots, v_n\}$, say **forward** k -hop neighbors of v_i
- From v_1 to v_n , in the i -th iteration, list maximal k -plexes with i as the minimum index of vertices.
 - Build **seed graph** $G_i = G[\{v_i\} \cup N_{\succ}(v_i) \cup N_{\succ}^2(v_i)]$.
 - Call BKPivot with given combinations of $N_{\succ}^2(v_i)$, say **seed set** S .
 - Validate maximality in G for k -plexes generated in G_i .

ListPlex

Part II




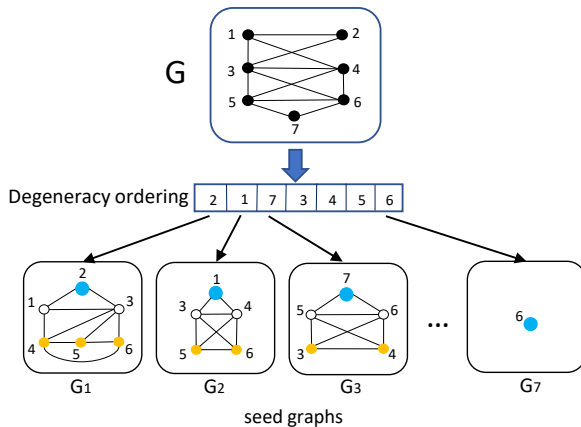
Figure 7: An example of ListPlex's Part II.

(L) sort V in degeneracy ordering η and induce seed graphs G_i for each $v_i \in \eta$.

(R) enumerate $S \subseteq N_{\mathcal{L}}^2(v_i)$ with bound $|S| \leq k - 1$ ($k = 3$) and call BKPIVOT with P_S, C_S, X_S .

ListPlex


Graph Decomposition



- Cohesive groups appear locally in large real-world graphs,
- Parallelism,
- Smaller scale and better locality,
- $D \ll \Delta$.

ListPlex

Seed Set



Call BKPivot
with Ps,Cs,Xs

S	Ps	Cs	Xs
\emptyset	{2}	{1,3}	{4,5,6}
{4}	{2,4}	{1,3}	{5,6}
{5}	{2,5}	{1,3}	{4,6}
{6}	{2,6}	{1,3}	{4,5}
{4,5}	{2,4,5}	{1,3}	{6}
{4,6}	{2,4,6}	{1,3}	{5}
{5,6}	{2,5,6}	{1,3}	{4}

- At most $k - 1$ vertices come from $N_{\prec}^2(v_i)$,
- Reducing candidates fast,

- $|N_{\prec}^2(v_i)|$ is potentially $D\Delta$,
- Pruning rules.

ListPlex

Listing Large Maximal k -Plexes

With lower bound l , a related lower bound l' can be derived.

Lemma 3

Assume $|P| \geq l$, for any vertex pair $u, v \in P$, $|N(u) \cap N(v) \cap P| \geq l'$

Removing unfruitful candidates from G_i reduces the scale of G_i .

- Consider vertex pair (v_i, u) , $u \in V_i$.

Dropping unfruitful seed sets S saves the forthcoming exponential search.

- Consider vertex pair (u, v) from some S .

1 Background

2 Algorithms

3 Implementation Techniques

4 Experiments

Reducing cache misses

Validating maximality in G for maximal k -plex of G_i suffers a high amount of cache misses.

Alleviation:

- Build a bipartite graph B_i for each G_i , which serves as a cache of currently useful data of large G ,
- Pruning bipartite graph just like seed graph.

Parallelization

ListPlex owns appealing parallel features.

How To:

- Parallelize searches of maximal k -plexes on each G_i , say T_i ,
- Split T_i when some cores are idle for better load balance,
- Construct degeneracy ordering and perform generated T_i in parallel.

- 1 Background
- 2 Algorithms
- 3 Implementation Techniques
- 4 Experiments**

Dataset

All graphs are taken from SNAP and LAW ⁴.

Table 1: Considered networks and their properties

Network	n	m	Δ	D
jazz	198	2742	100	29
ca-grqc	5241	14484	81	43
gnutella08	6301	41554	97	10
wiki-vote	7116	100763	1065	53
lastfm	7624	55612	216	20
as-caida	26475	53381	2628	22
soc-epinions	75888	405739	3044	67
soc-slashdot	82144	500480	2548	54
email-euall	265214	365569	7636	37
amazon0505	410236	2439436	2760	10
in-2004	1353703	13126172	21869	488
soc-pokec	1632803	22301964	14854	47
as-skitter	1696415	11095298	35455	111
soc-livejournal	4847571	68993773	14815	360
arabic-2005	22744080	639999458	575628	3247
uk-2005	39459925	936364282	1372171	584
it-2004	41291594	1150725436	1243927	3209
webbase-2001	118142155	1019903190	816127	1506

⁴<http://law.di.unimi.it/>

Listing All Maximal k -Plexes

Table 2: Listing all maximal k -plexes in small graphs

Network	k	# k -plexes	The running time (s)				Speedup
			BKPlex	BKPivot	ListPlex	ListPlex(16)	
jazz	2	35214	648.864	0.29	0.086	0.408	0.211
jazz	3	3602575	772.826	17.55	6.477	0.832	7.785
jazz	4	193056583	3226.746	829.40	417.646	26.187	15.949
ca-grqc	2	13718439	OOT	1858.02	649.985	40.880	15.899
gnutella08	2	19866959	1500.208	3627.57	1117.858	70.207	15.922
wiki-vote	2	66193264	10356.553	10671.92	1526.884	95.656	15.962
lastfm	2	29086855	2643.394	6676.89	1989.701	124.525	15.978

Listing Large Maximal k -Plexes

Table 3: The running time of listing large maximal k -plexes from small and medium graphs by CommuPlex⁵, D2K⁶ and ListPlex.

Graph ($ V , E $)	k	l	# k -plexes	The running time (s)			Graph ($ V , E $)	k	l	# k -plexes	The running time (s)		
				Commuplex	D2K	ListPlex					Commuplex	D2K	ListPlex
jazz (198, 2742)	4	12	2745953	25.218	33.054	4.498	wiki-vote (7116, 100763)	2	12	2919931	75.871	115.757	17.653
lastfm (7624, 55612)	4	12	1827337	20.724	23.991	4.586		20	52	4.52	11.289	0.591	
as-caida (26475, 53381)	3	12	281251	5.684	13.421	0.867		30	0	1.033	0.027	0.091	
amazon0505 (410236, 2439436)	4	12	15939891	300.388	785.506	47.98		12	458153397	OOT	OOT	2185.598	
email-euall (265214, 365569)	2	12	376	1.825	0.641	0.137		3	20	156727	595.636	1852.186	9.384
	3	12	6347	11.359	0.77	0.286		30	0	1.072	0.029	0.1	
	4	12	105649	47.049	5.338	1.171		20	46729532	OOT	OOT	1174.2	
	2	12	412779	8.793	11.199	1.946		30	0	9.17	3.627	0.112	
soc-slashdot (82144, 500480)	3	12	32639016	619.384	1043.266	91.62		2	12	7679906	1537.506	172.987	47.475
	20	2637	10.754	53.691	0.429			20	94184	1064.371	20.03	15.161	
	4	20	1707177	825.126	3800.889	24.089		30	3	662.64	8.637	9.557	
	12	27208777	376.071	213.141	59.42		12	520888893	OOT	OOT	1607.285		
as-skitter (1696415, 11095298)	2	20	11411028	227.016	137.159	32.988	3	20	5911456	1470.536	856.393	46.262	
	30	453	10.77	16.481	0.688		30	5	717.425	9.993	10.127		
	4	30	502699966	OOT	OOT	6680.261	4	20	318035938	34048.155	OOT	1825.216	
	2	50	47969775	OOT	OOT	520.884	30	4515	1140.117	111.987	11.211		
in-2004 (1353703, 13126172)	2	100	0	1.793	2.951	0.716	soc-epinions (75888, 405739)	2	12	49823056	843.9	735.589	193.307
	3	50	21070497438	OOT	OOT	OOT		20	3322167	137.427	180.061	19.382	
	3	100	0	2.37	3.285	0.718		30	0	8.995	12.109	0.492	
	2	50	25855779	7663.843	576.06	150.212		3	20	548634119	27037.614	35525.693	3072.267
com-livejournal (4847571, 68993773)	2	100	9978037	5899.638	256.225	72.063		4	30	13172906	OOT	OOT	661.103
	3	50	29045783792	OOT	OOT	OOT		2	340	650322	2284.435	OOT	109.382
	3	100	4257410159	OOT	OOT	28384.76		2	345	0	57.548	13589.487	6.914
	3	340	555718694	OOT	OOT	22863.467		3	345	3963139	24861.871	OOT	826.183

⁵Zhou et al. (2020)

⁶Conte et al. (2018)

Listing Large Maximal k -Plexes

Table 4: The parallel running time of large networks by ListPlex and D2K with 16 threads.

Graph ($ V , E $)	k	l	# k -plexes	The running time (s)	
				D2K(16)	ListPlex(16)
arabic-2005 (22744080, 639999458)	2	800	224870903	2195.272	714.159
	2	1000	236897	151.328	40.202
	3	800	>25062182205	OOT	OOT
	3	1000	34155502	587.967	128.737
uk-2005 (39459925, 936364282)	2	250	106243475	OOT	355.855
	2	500	256406	318.118	35.001
	3	250	>18336111409	OOT	OOT
	3	500	28199814	9506.661	121.726
it-2004 (41291594, 1150725436)	2	2000	675111	340.904	41.983
	2	3000	675111	307.735	38.468
	3	2000	197679229	4254.456	724.979
	3	3000	197679229	4235.389	715.002
webbase-2001 (118142155, 1019903190)	2	800	1599005	374.134	54.19
	2	1000	1164383	346.393	53.651
	3	800	1785341050	36116.817	5521.386
	3	1000	1484341137	35005.343	6960.816




Figure 8: The speedup of ListPlex for the large graphs with different parameters.

Excluding unfruitful seed sets

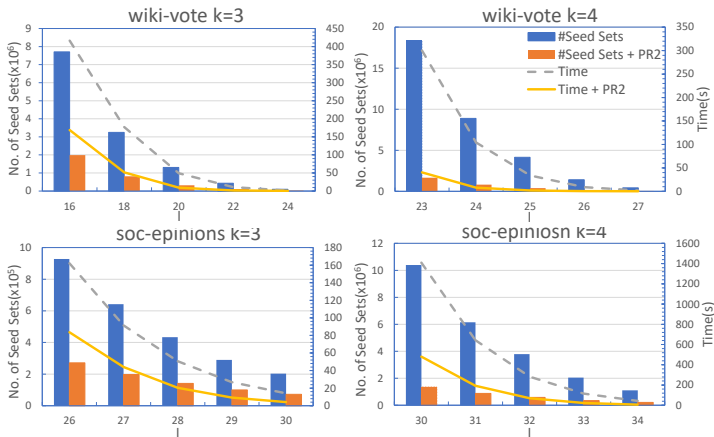


Figure 9: The number of seed sets and running time with and without Prune Rule 2.

Reducing cache misses




Figure 10: The total number of data cache misses and the running time with and without using bipartite graph B_i .

Q&A

Thanks!